

Chapter-4

Naming

Introduction

- names play an important role to:
 - share resources
 - uniquely identify entities
 - refer to locations
 - etc.
- an important issue is that a name can be **resolved** to the entity it refers
- to resolve names, it is necessary to implement a **naming system**
- in a distributed system, the implementation of a naming system is itself often distributed, unlike in non-distributed systems
- **efficiency** and **scalability** of the naming system are the main issues

1. Naming Entities

■ Names, Identifiers, and Addresses

- a **name** in a distributed system is a string of bits or characters that is used to refer to an **entity**
- an **entity** is anything; e.g., resources such as hosts, printers, disks, files, objects, processes, users, ...
- entities can be **operated** on; e.g., a resource such as a printer offers an interface containing operations for printing a document, requesting the status of a job, ...
- to operate on an entity, it is necessary to access it through its **access point**, itself an entity (special)

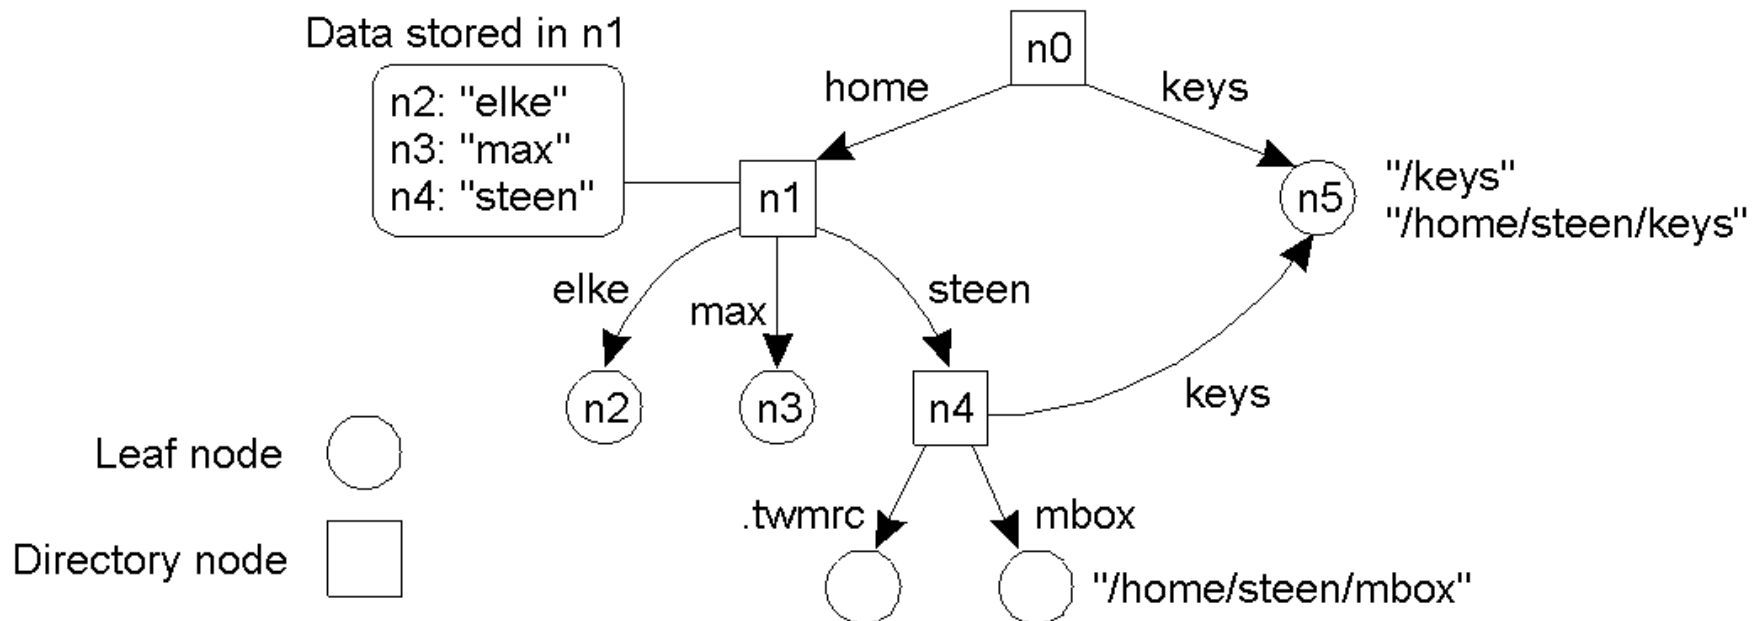
- access point
 - the name of an access point is called an address (such as IP address and port number as used by the transport layer)
 - the address of the access point of an entity is also referred to as the address of the entity
 - an entity can have more than one access point (similar to accessing an individual through different telephone numbers)
 - an entity may change its access point in the course of time (e.g., a mobile computer getting a new IP address as it moves)

- an **address** is a special kind of name
 - it refers to at most one entity
 - each entity is referred by at most one address; even when replicated such as in Web pages
 - an entity may change an access point, or an access point may be reassigned to a different entity (like telephone numbers in offices)
 - separating the name of an entity and its address makes it easier and more flexible; such a name is called **location independent**
- there are also other types of names that uniquely identify an entity; in any case an **identifier** is a name with the following properties
 - it refers to at most one entity
 - each entity is referred by at most one identifier
 - it always refers to the same entity (never reused)
- identifiers allow us to unambiguously refer to an entity

- examples
 - name of an FTP server (entity)
 - URL of the FTP server
 - address of the FTP server
 - IP number: port number
 - the address of the FTP server may change

2. Name Spaces and Name Resolution

- names in a distributed system are organized into a **name space**
- a name space is generally organized as a labeled, directed graph with two types of nodes
 - leaf node**: represents the named entity and stores information such as its address or the state of that entity
 - directory node**: a special entity that has a number of outgoing edges, each labeled with a name



a general naming graph with a single root node

- each node in a naming graph is considered as another entity with an identifier
- a directory node stores a table in which an outgoing edge is represented as a pair (*edge label*, *node identifier*), called a **directory table**
- each path in a naming graph can be referred by the sequence of labels corresponding to the edges of the path and the first node in the path, such as
 $N:\langle label1, label2, \dots, labeln \rangle$, where N refers to the first node in the path
- such a sequence is called a **path name**
- if the first node is the root of the naming graph, it is called an **absolute path name**; otherwise it is a **relative path name**
- instead of the path name $n0:\langle home, steen, mbox \rangle$, we often use its string representation `/home/steen/mbox`
- there may also be several paths leading to the same node, e.g., node $n5$ can be represented as `/keys` or `/home/steen/keys`

- although the above naming graph is **directed acyclic graph** (a node can have more than one incoming edge but is not permitted to have a cycle), the common way is to use a tree (hierarchical) with a single root (as is used in file systems)
 - in a tree structure, each node except the root has exactly one incoming edge; the root has no incoming edges
 - ➡ each node also has exactly one associated (absolute) path name

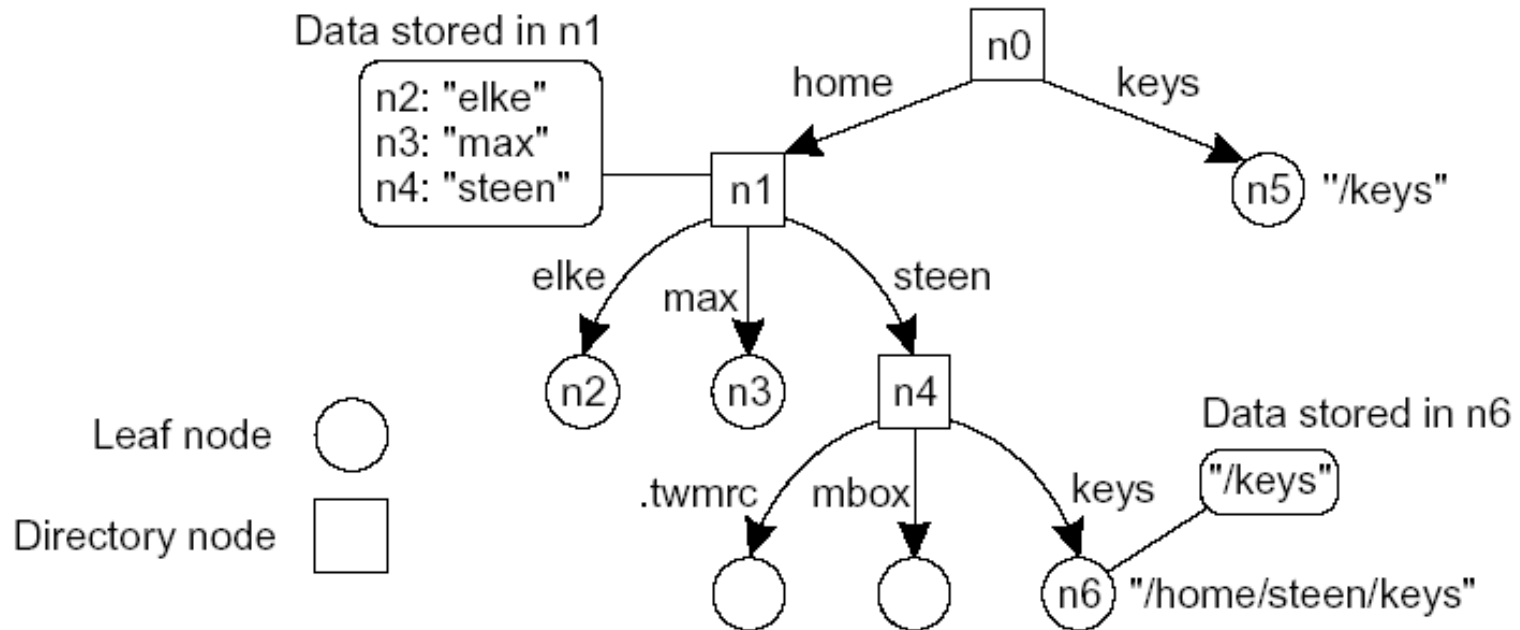
■ Name Resolution

- given a path name, the process of looking up a name stored in the node is referred to as **name resolution**; it consists of finding the address when the name is given (by following the path)

■ Linking and Mounting

- **Linking**: giving another name for the same entity (an **alias**)
e.g., environment variables in UNIX such as HOME that refer to the home directory of a user
- two types of links (or two ways to implement an alias):
 - **hard link**: to allow multiple absolute path names to refer to the same node in a naming graph
e.g., in the previous graph, there are two different path names for node n5: */keys* and */home/steen/keys*

- **symbolic link**: representing an entity by a leaf node and instead of storing the address or state of the entity, the node stores an absolute path name



the concept of a symbolic link explained in a naming graph

- when first resolving an absolute path name stored in a node (e.g., /home/steen/keys in node n6), name resolution will return the path name stored in the node (/keys), at which point it can continue with resolving that new path name

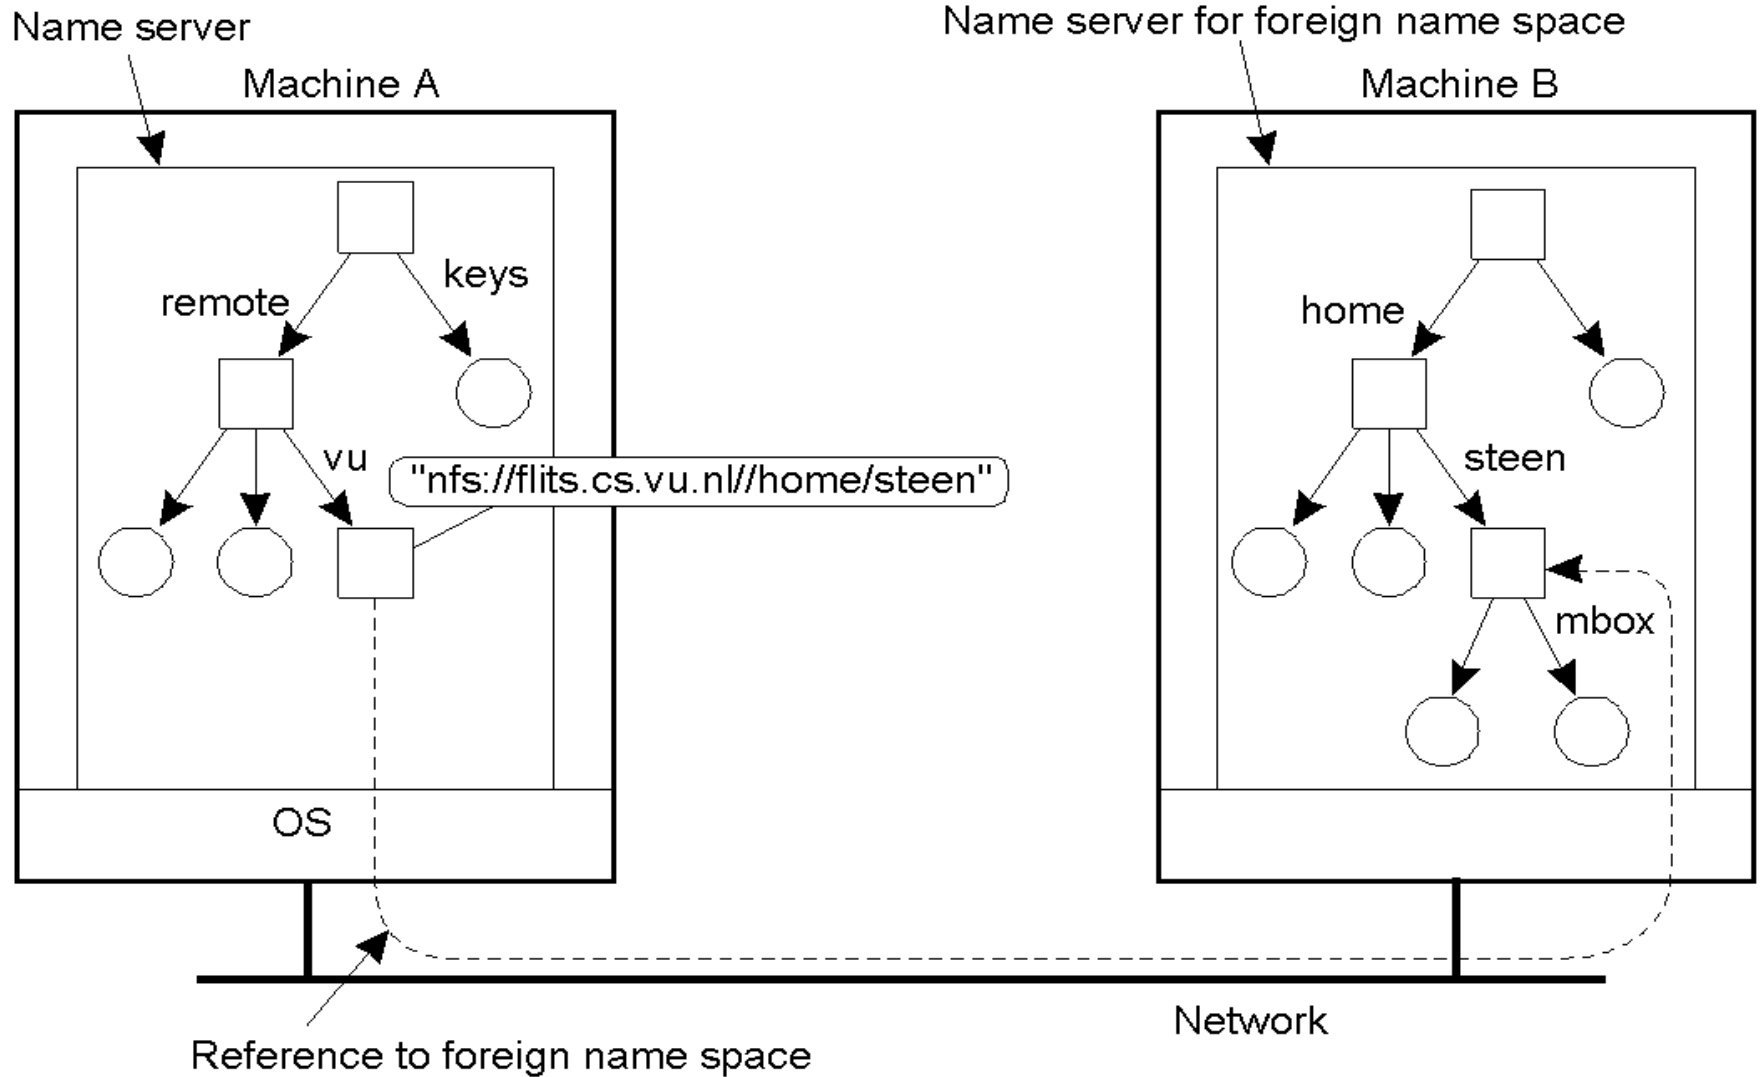
- so far name resolution was discussed as taking place within a single name space
- name resolution can also be used to merge different name spaces in a transparent way
- two methods: **mounting** and **adding a new root node and making the existing root nodes its children**

1. Mounting

- as an example, consider a mounted file system, which can be generalized to other name spaces as well
- let a directory node store the directory node from a different (foreign) name space
- the directory node storing the node identifier is called a **mount point**
- the directory node in the foreign name space is called a **mounting point**, normally the root of a name space
- during name resolution, the mounting point is looked up and resolution proceeds by accessing its directory table

- consider a collection of name spaces distributed across different machines (each name space implemented by a different server)
- to mount a foreign name space in a DS, the following are at least required
 - the name of an **access protocol** (for communication)
 - the name of the **server**
 - the name of the **mounting point**
- each of these names needs to be resolved
 - to the implementation of the **protocol**
 - to an **address** where the server can be reached
 - to a **node identifier** in the foreign name space
- the three names can be listed as a URL

- example: Sun's Network File System (NFS) is a distributed file system with a protocol that describes how a client can access a file stored on a (remote) NFS file server
 - an NFS URL may look like [nfs://flits.cs.vu.nl/home/steen](#)
 - [nfs](#) is an implementation of a protocol
 - [flits.cs.vu.nl](#) is a server name to be resolved using DNS
 - [/home/steen](#) is resolved by the server
 - e.g., the subdirectory [/remote](#) includes mount points for foreign name spaces on the client machine
 - a directory node named [/remote/vu](#) is used to store [nfs://flits.cs.vu.nl/home/steen](#)
 - consider [/remote/vu/mbox](#)
 - this name is resolved by starting at the root directory on the client's machine until node [/remote/vu](#), which returns the URL [nfs://flits.cs.vu.nl/home/steen](#)
 - this leads the client machine to contact [flits.cs.vu.nl](#) using the NFS protocol
 - then the file [mbox](#) is read in the directory [/home/steen](#)

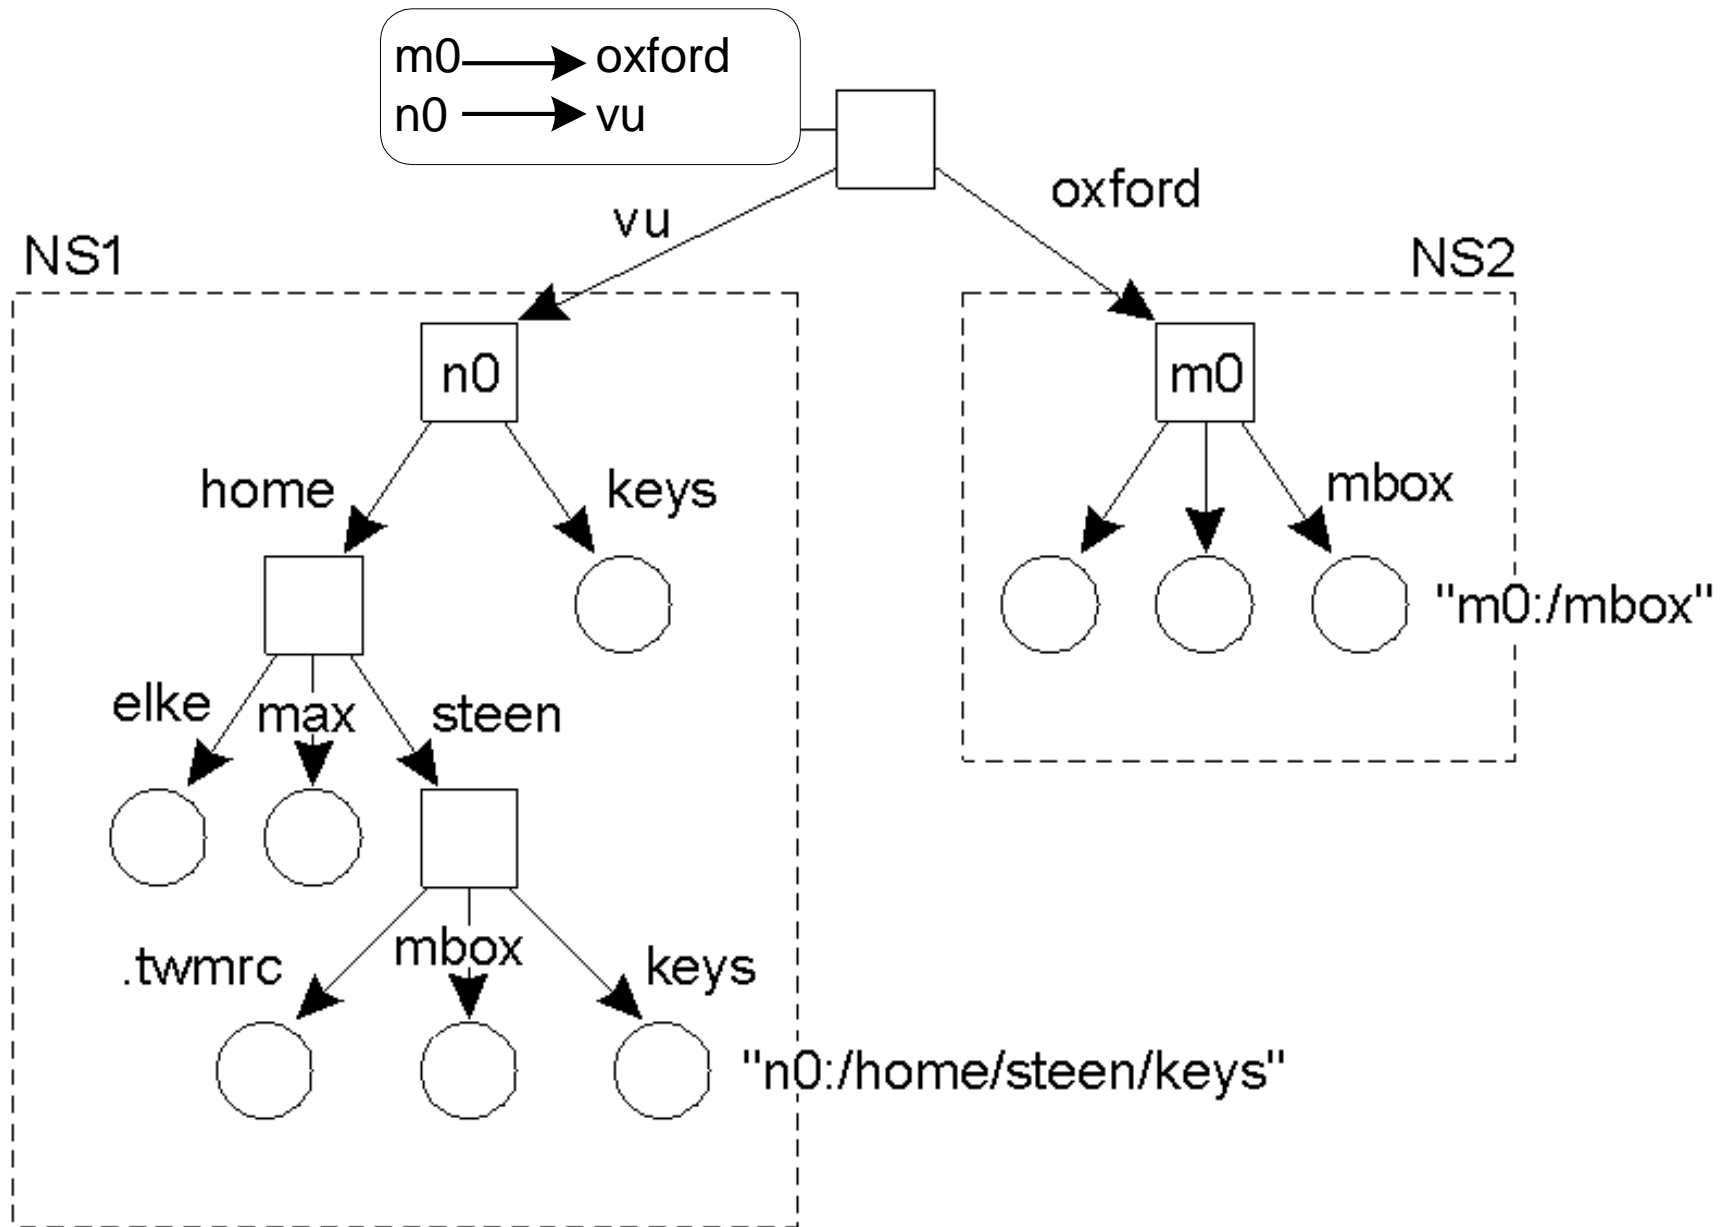


mounting remote name spaces through a specific process protocol

- distributed systems that allow mounting a remote file system also allow to execute some commands
- example commands to access the file system
 `cd /remote/vu`
- by doing so the user is not supposed to worry about the details of the actual access; the name space on the local machine and that on the remote machine look to form a single name space

2. Add a new root node and make the existing root nodes its children

- a method followed in GNS (Global Name Service by DEC)
- problem: existing names need to be changed
e.g., the absolute path name `/home/steen` has now changed to a relative path name and corresponds to the absolute path name `/vu/home/steen`
- hence the system must expand `no:/home/steen` to `/vu/home/steen` without the awareness of users
- this requires storing a mapping table (with entries such as `n0→vu`) when a new root node is added
- merging thousands of name spaces may lead to performance problems



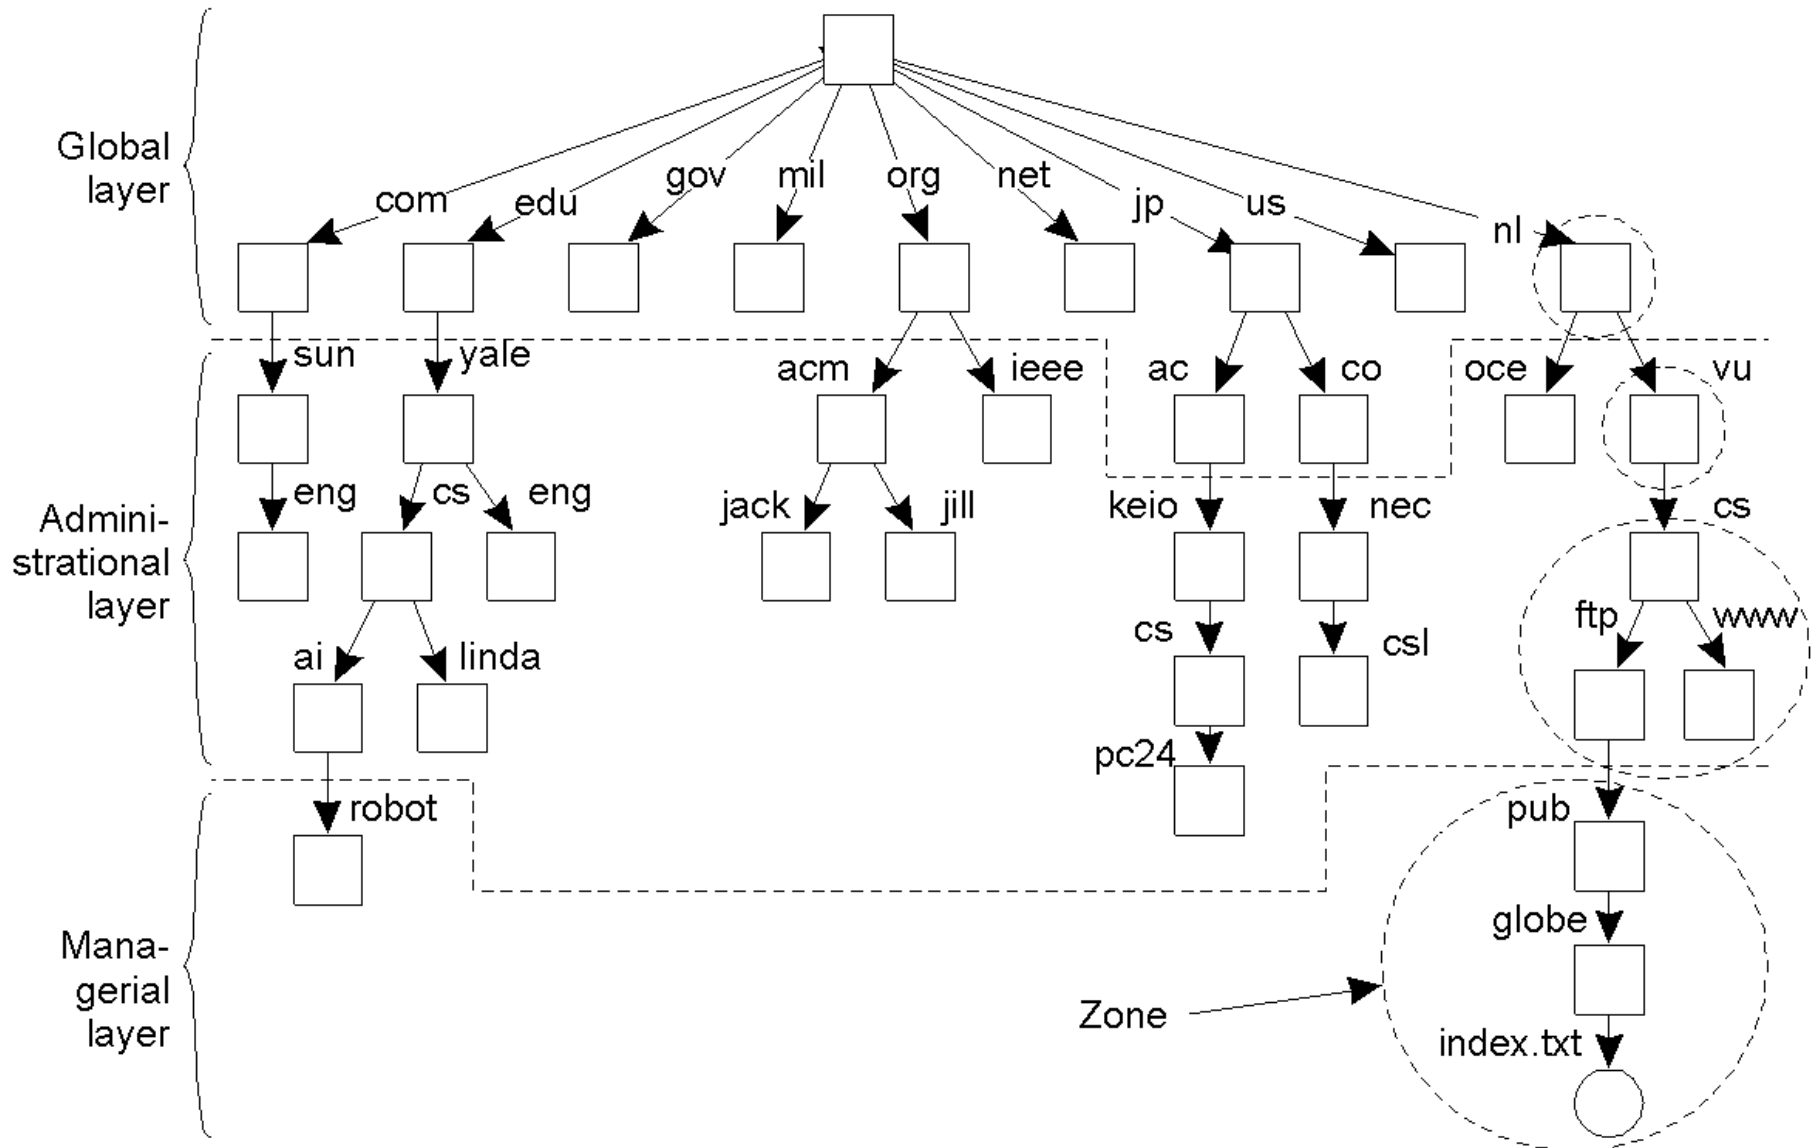
■ The Implementation of a Name Space

- a name space forms the heart of a naming service
- a naming service allows users and processes to **add**, **remove**, and **lookup** names
- a naming service is implemented by **name servers**
- for a distributed system on a single LAN, a single server might suffice; for a large-scale distributed system the implementation of a name space is distributed over multiple name servers

■ Name Space Distribution

- in large scale distributed systems, it is necessary to distribute the name service over multiple name servers, usually organized hierarchically
- a name service can be partitioned into logical layers
- the following three layers can be distinguished (Cheriton and Mann)

- **global layer**
 - formed by highest level nodes (root node and nodes close to it or its children)
 - nodes on this layer are characterized by their **stability**, i.e., directory tables are rarely changed
 - they may represent organizations, groups of organizations, ..., where names are stored in the name space
- **administrational layer**
 - groups of entities that belong to the same organization or administrative unit, e.g., departments
 - relatively stable
- **managerial layer**
 - nodes that may change regularly, e.g., nodes representing hosts of a LAN, shared files such as libraries or binaries, ...
 - nodes are managed not only by system administrators, but also by end users



an example partitioning of the DNS name space, including Internet-accessible files, into three layers

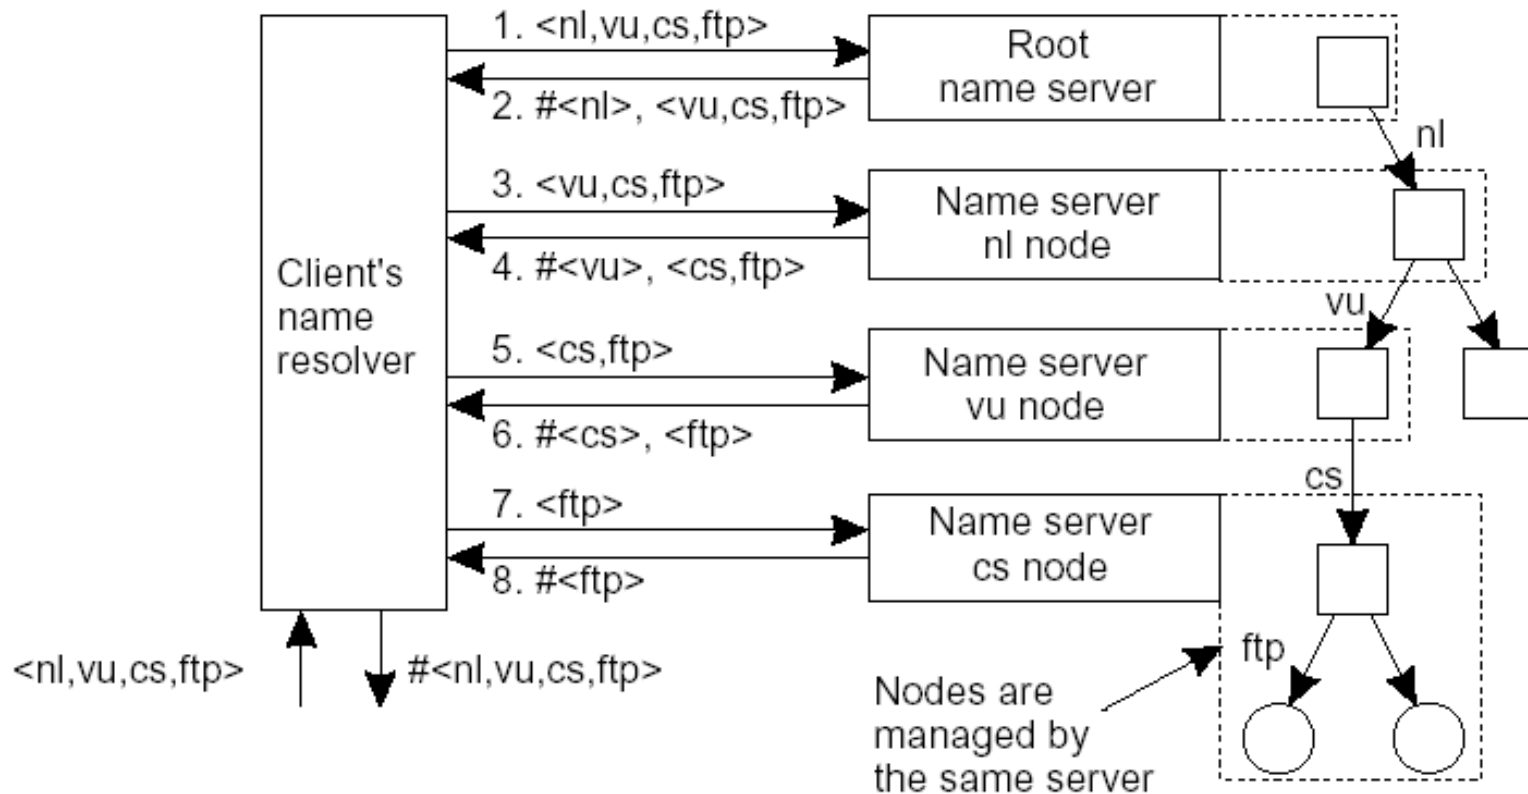
- the name space is divided into nonoverlapping parts, called **zones** in DNS
- a zone is a part of the name space that is implemented by a separate **name server**
- some requirements of servers at different layers
 - **performance** (responsiveness to lookups), **availability** (failure rate), etc.
 - high **availability** is critical for the **global layer**, since name resolution cannot proceed beyond the failing server; it is also important at the **administrational layer** for clients in the same organization
 - **performance** is very important in the **lowest layer**, since results of lookups can be cached and used due to the relative stability of the higher layers
 - they may be enhanced by **client side caching** (global and administrational layers since names do not change often) and **replication**; they create implementation problems since they may introduce inconsistency problems.

■ Implementation of Name Resolution

- recall that name resolution consists of finding the **address** when the **name** is given
- assume that name servers are not replicated and that no client-side caches are allowed
- each client has access to a local **name resolver**, responsible for ensuring that the name resolution process is carried out
- e.g., assume the path name
 `root:<nl, vu, cs, ftp, pub, globe, index.txt>`
 is to be resolved
 or using a URL notation, this path name would correspond to
 `ftp://ftp.cs.vu.nl/pub/globe/index.txt`
- two ways of implementing name resolution
 - **iterative name resolution**
 - **recursive name resolution**

■ Iterative

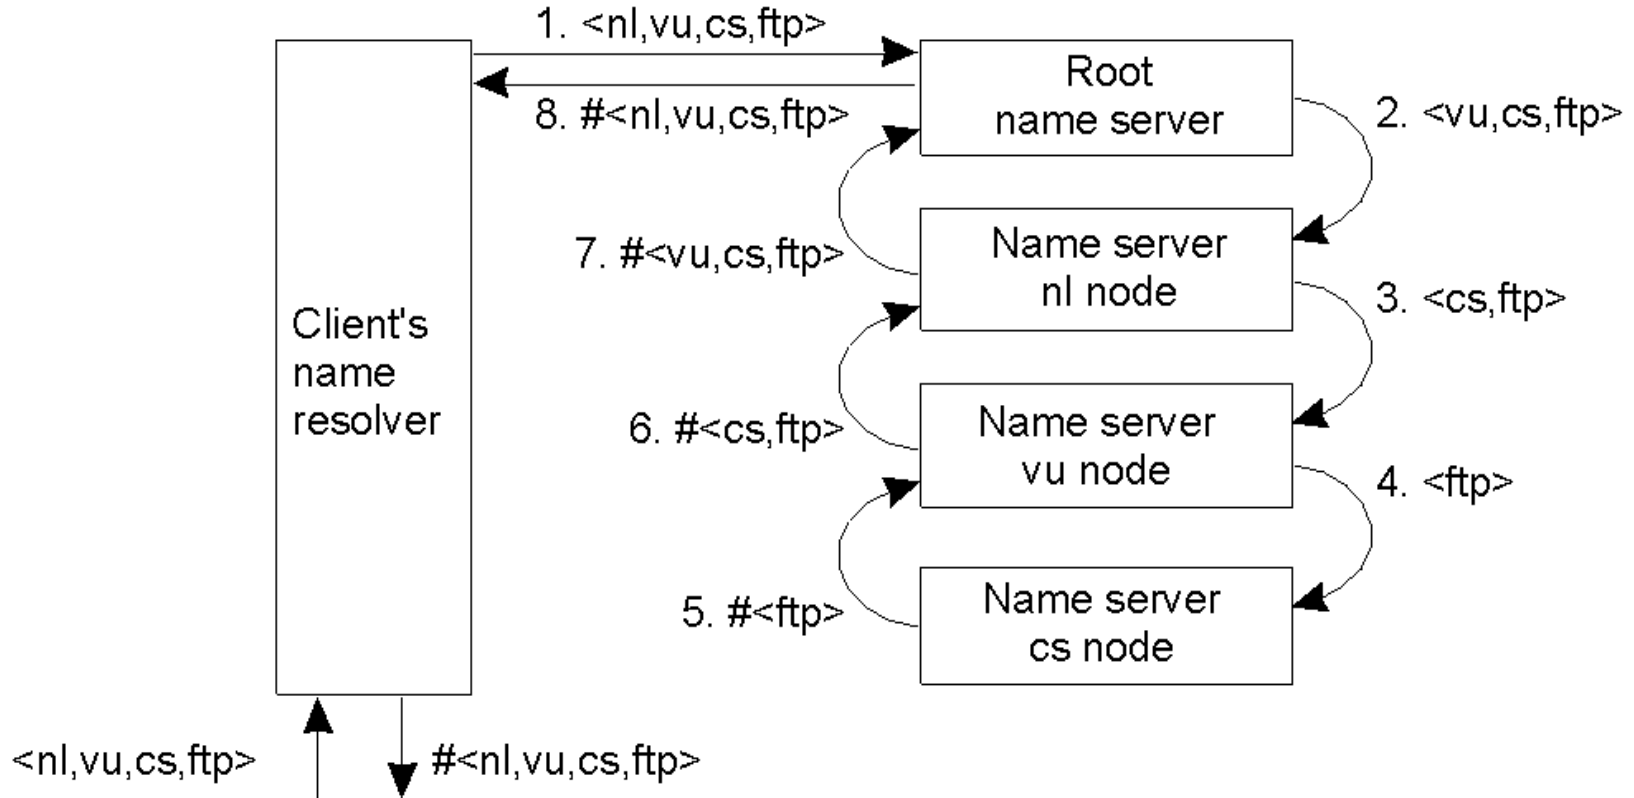
- a name resolver hands over the complete name to the root name server
- the root server will resolve the name as far as it can and return the result to the client; at the minimum it can resolve the first level and sends the name of the first level name server to the client
- the client calls the first level name server, then the second, ..., until it finds the address of the entity



the principle of iterative name resolution

■ Recursive

- a name resolver hands over the whole name to the root name server
- the root server will try to resolve the name and if it can't, it requests the first level name server to resolve it and to return the address
- the first level will do the same thing recursively



the principle of recursive name resolution

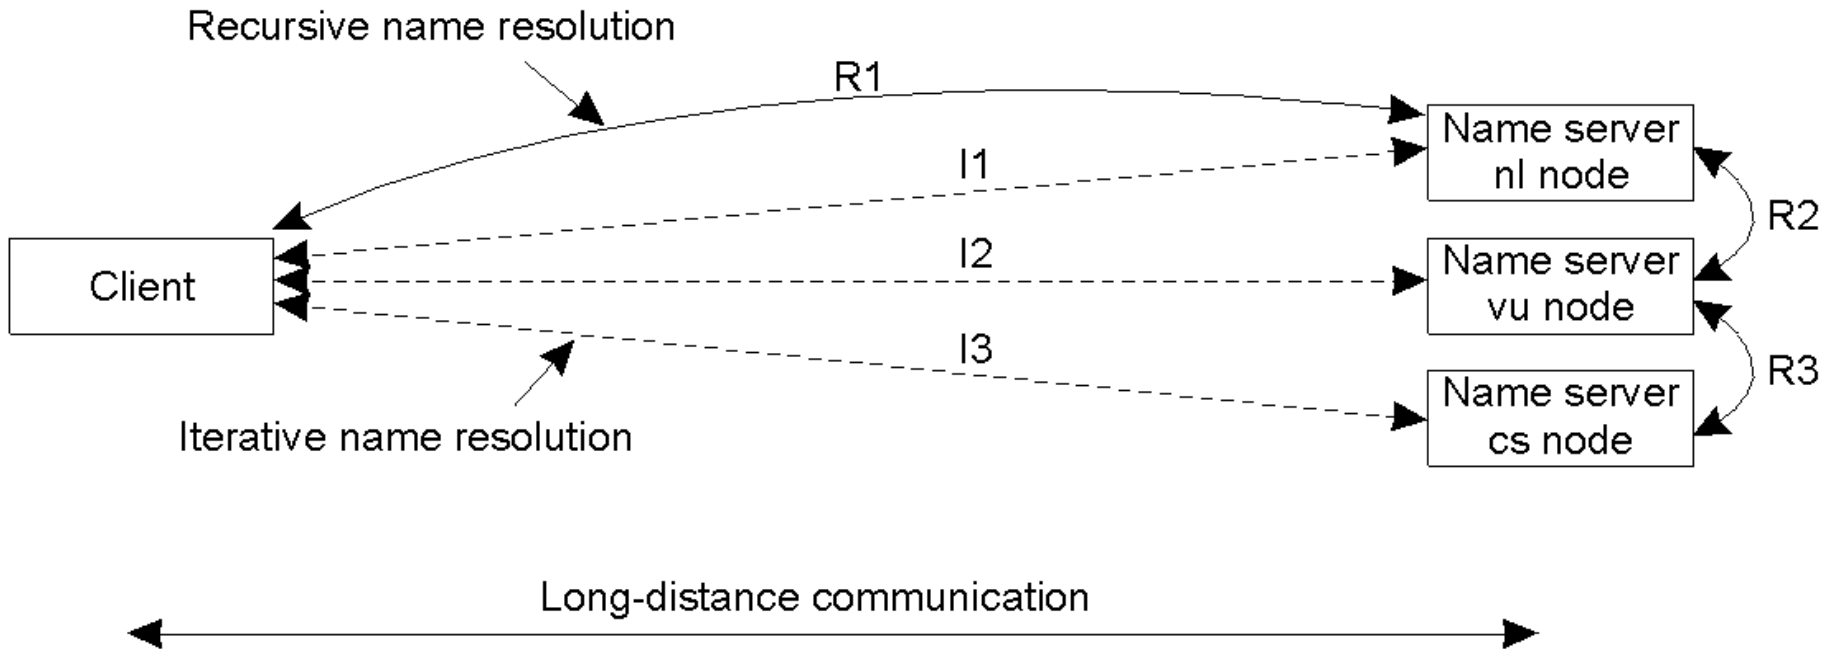
- Advantages and drawbacks

- recursive name resolution puts a higher performance demand on each name server; hence name servers in the global layer support only iterative name resolution
- caching is more effective with recursive name resolution; each name server gradually learns the address of each name server responsible for implementing lower-level nodes; eventually lookup operations can be handled efficiently

Server for node	Should resolve	Looks up	Passes to child	Receives and caches	Returns to requester
cs	<ftp>	#<ftp>	--	--	#<ftp>
vu	<cs,ftp>	#<cs>	<ftp>	#<ftp>	#<cs> #<cs, ftp>
nl	<vu,cs,ftp>	#<vu>	<cs,ftp>	#<cs> #<cs,ftp>	#<vu> #<vu,cs> #<vu,cs,ftp>
root	<nl,vu,cs,ftp>	#<nl>	<vu,cs,ftp>	#<vu> #<vu,cs> #<vu,cs,ftp>	#<nl> #<nl,vu> #<nl,vu,cs> #<nl,vu,cs,ftp>

recursive name resolution of <nl, vu, cs, ftp>; name servers cache intermediate results for subsequent lookups

- communication costs may be reduced in recursive name resolution



the comparison between recursive and iterative name resolution with respect to communication costs; assume the client is in Ethiopia and the name servers in the Netherlands

- Summary

Method	Advantage(s)
Recursive	Less Communication cost; Caching is more effective
Iterative	Less performance demand on name servers

- Example 1 - The Domain Name System (DNS)

- one of the largest distributed naming services is the Internet DNS
- it is used for looking up host addresses and mail servers
- hierarchical, defined in an inverted tree structure with the root at the top
- the tree can have only 128 levels

